

Diamonds ML R Notebook

Diamonds ML R Notebook Robert M. Taylor, PhD

This notebook is to demonstrate Exploratory Data Analysis (EDA), visualizations, and machine learning in R on the diamonds dataset that is available in R. “Price” will be our target.s

I'll first import the libraries I'll use.

```
library(ggplot2)
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2
```

```
library(rpart)
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

```
library(tidyr)
library(modelr)
```

Load the dataset

```
data(diamonds)
```

I'll just look at/inspect the dataset first. This is a clean data set so data cleaning will not be needed or demonstrated in this notebook.

```
head(diamonds)
```

```
## # A tibble: 6 x 10
##   carat cut      color clarity depth table price     x     y     z
##   <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
## 1 0.23 Ideal     E     SI2     61.5    55   326   3.95  3.98  2.43
## 2 0.21 Premium  E     SI1     59.8    61   326   3.89  3.84  2.31
## 3 0.23 Good     E     VS1     56.9    65   327   4.05  4.07  2.31
## 4 0.290 Premium  I     VS2     62.4    58   334   4.2   4.23  2.63
## 5 0.31 Good     J     SI2     63.3    58   335   4.34  4.35  2.75
## 6 0.24 Very Good J     VVS2     62.8    57   336   3.94  3.96  2.48
```

What are the dimensions of the dataset?

```
dim(diamonds)
```

```
## [1] 53940      10
```

So there are 53,940 rows and 10 feature columns.

I'll now get 1) a summary and 2) the structure of the data...

```
summary(diamonds)
```

```
##      carat      cut      color      clarity      depth
##  Min.   :0.2000   Fair      : 1610   D: 6775   SI1      :13065   Min.   :43.00
##  1st Qu.:0.4000   Good      : 4906   E: 9797   VS2      :12258   1st Qu.:61.00
##  Median :0.7000   Very Good:12082   F: 9542   SI2      : 9194   Median :61.80
##  Mean   :0.7979   Premium  :13791   G:11292   VS1      : 8171   Mean   :61.75
##  3rd Qu.:1.0400   Ideal    :21551   H: 8304   VVS2     : 5066   3rd Qu.:62.50
##  Max.   :5.0100                I: 5422   VVS1     : 3655   Max.   :79.00
##                                J: 2808   (Other): 2531
##      table      price      x      y
##  Min.   :43.00   Min.   : 326   Min.   : 0.000   Min.   : 0.000
##  1st Qu.:56.00   1st Qu.: 950   1st Qu.: 4.710   1st Qu.: 4.720
##  Median :57.00   Median : 2401   Median : 5.700   Median : 5.710
##  Mean   :57.46   Mean   : 3933   Mean   : 5.731   Mean   : 5.735
##  3rd Qu.:59.00   3rd Qu.: 5324   3rd Qu.: 6.540   3rd Qu.: 6.540
##  Max.   :95.00   Max.   :18823   Max.   :10.740   Max.   :58.900
##
##      z
##  Min.   : 0.000
##  1st Qu.: 2.910
##  Median : 3.530
##  Mean   : 3.539
##  3rd Qu.: 4.040
##  Max.   :31.800
##
```

I see that there is an (Other) variable for Clarity. I want to look at that closer.

```
unique(diamonds$clarity)
```

```
## [1] SI2 SI1 VS1 VS2 VVS2 VVS1 I1 IF
## Levels: I1 < SI2 < SI1 < VS2 < VS1 < VVS2 < VVS1 < IF
```

So, everything appears good. The summary has just grouped the I1 and IF clarities in the count values in the summary table above.

We also see that the price ranges from a minimum price of \$326 to a max of \$18,823

I'll go ahead and look at the structure of the dataset...

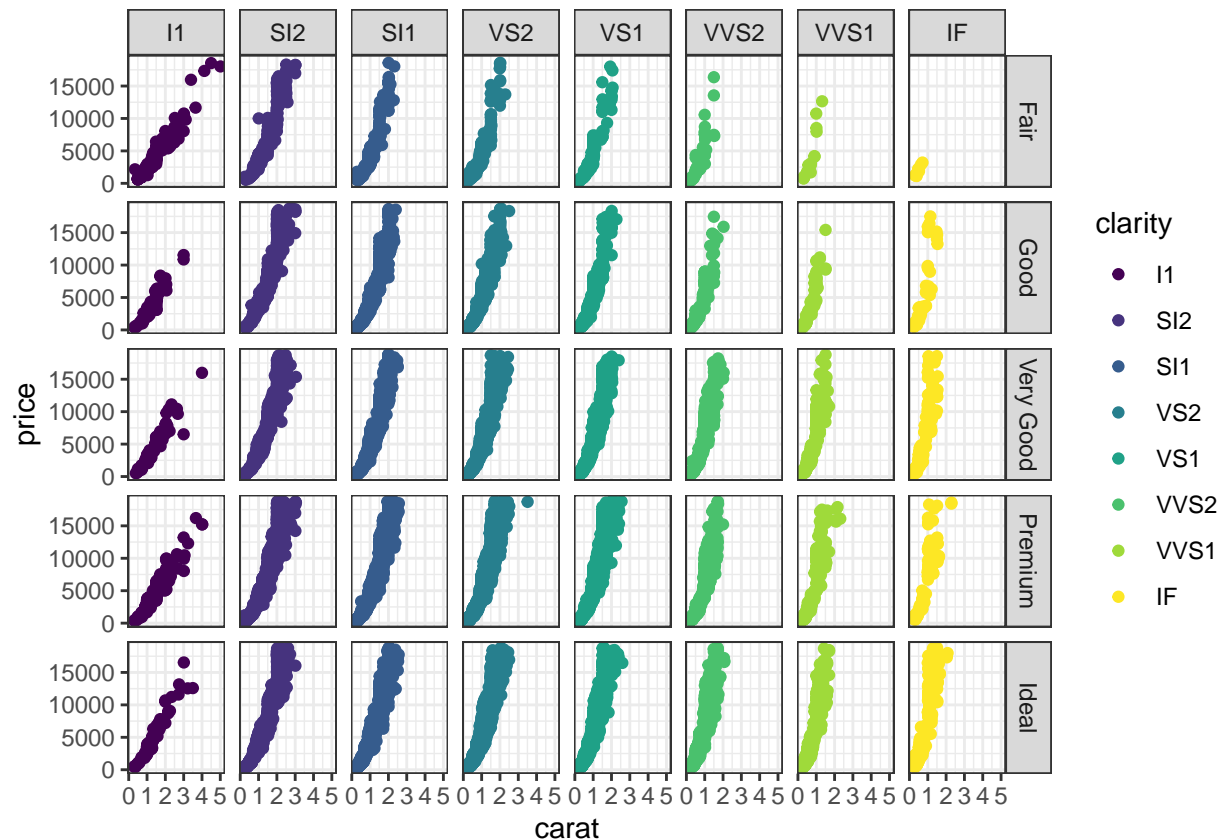
```
str(diamonds)
```

```
## tibble [53,940 x 10] (S3: tbl_df/tbl/data.frame)
##  $ carat   : num [1:53940] 0.23 0.21 0.23 0.29 0.31 0.24 0.24 0.26 0.22 0.23 ...
##  $ cut     : Ord.factor w/ 5 levels "Fair"<"Good"<...: 5 4 2 4 2 3 3 3 1 3 ...
##  $ color   : Ord.factor w/ 7 levels "D"<"E"<"F"<"G"<...: 2 2 2 6 7 7 6 5 2 5 ...
##  $ clarity : Ord.factor w/ 8 levels "I1"<"SI2"<"SI1"<...: 2 3 5 4 2 6 7 3 4 5 ...
##  $ depth   : num [1:53940] 61.5 59.8 56.9 62.4 63.3 62.8 62.3 61.9 65.1 59.4 ...
##  $ table   : num [1:53940] 55 61 65 58 58 57 57 55 61 61 ...
##  $ price   : int [1:53940] 326 326 327 334 335 336 336 337 337 338 ...
##  $ x       : num [1:53940] 3.95 3.89 4.05 4.2 4.34 3.94 3.95 4.07 3.87 4 ...
##  $ y       : num [1:53940] 3.98 3.84 4.07 4.23 4.35 3.96 3.98 4.11 3.78 4.05 ...
```

```
## $ z      : num [1:53940] 2.43 2.31 2.31 2.63 2.75 2.48 2.47 2.53 2.49 2.39 ...
```


I'll now look at carat vs price using ggplot2

```
g <- ggplot(diamonds, aes(x=carat, y=price))
g +
  geom_point(aes(color=clarity)) +
  facet_grid(cut ~ clarity) +
  theme_bw()
```




I'll use a boxplot to look at cut vs. price.

```
diamonds$cut = as.factor(diamonds$cut)
g <- ggplot(diamonds, aes(x=cut, y=price))
g +
  geom_boxplot() +
  facet_grid(~clarity) +
  theme(axis.text.x = element_text(angle = 90, face = "bold", color = "#993333",
    size = 9)) +
  theme(axis.text.y = element_text(face = "bold", color = "#993333"))
```




We can also look at the carat vs. color vs. price

```
g <- ggplot(diamonds, aes(x=carat, y=price))
g +
  geom_point(aes(color = color))
```




now look at the distribution of carat and price.

```
g <- ggplot(diamonds)
g +  
  geom_density(aes(x=carat), fill="gray50")
```



```
g <- ggplot(diamonds)
g +  
  geom_density(aes(x=price), fill="gray")
```



So, thus far, we can now see that most diamonds are < 2 carats and < ~\$2500 (although a second peak can be seen around \$4000).

M.L.

DECISION TREE BUILD MODEL I'll first use a decision tree model to predict the diamond prices.

```
colnames(diamonds)
```

```
## [1] "carat" "cut" "color" "clarity" "depth" "table" "price"
## [8] "x" "y" "z"
```

SPLIT THE DATA (70/30 split)

```
splitData <- resample_partition(diamonds, c(test=0.3, train=0.7))
```

How many cases are in the test and training sets?

```
lapply(splitData, dim)
```


```
## $test
## [1] 16181 10
##
## $train
## [1] 37759 10
```

```
# train the tree
```

```
fit <- rpart(price ~ carat + cut + color + clarity + depth + table + x + y + z, data = splitData$train)
```

```
#plot the regression tree
```

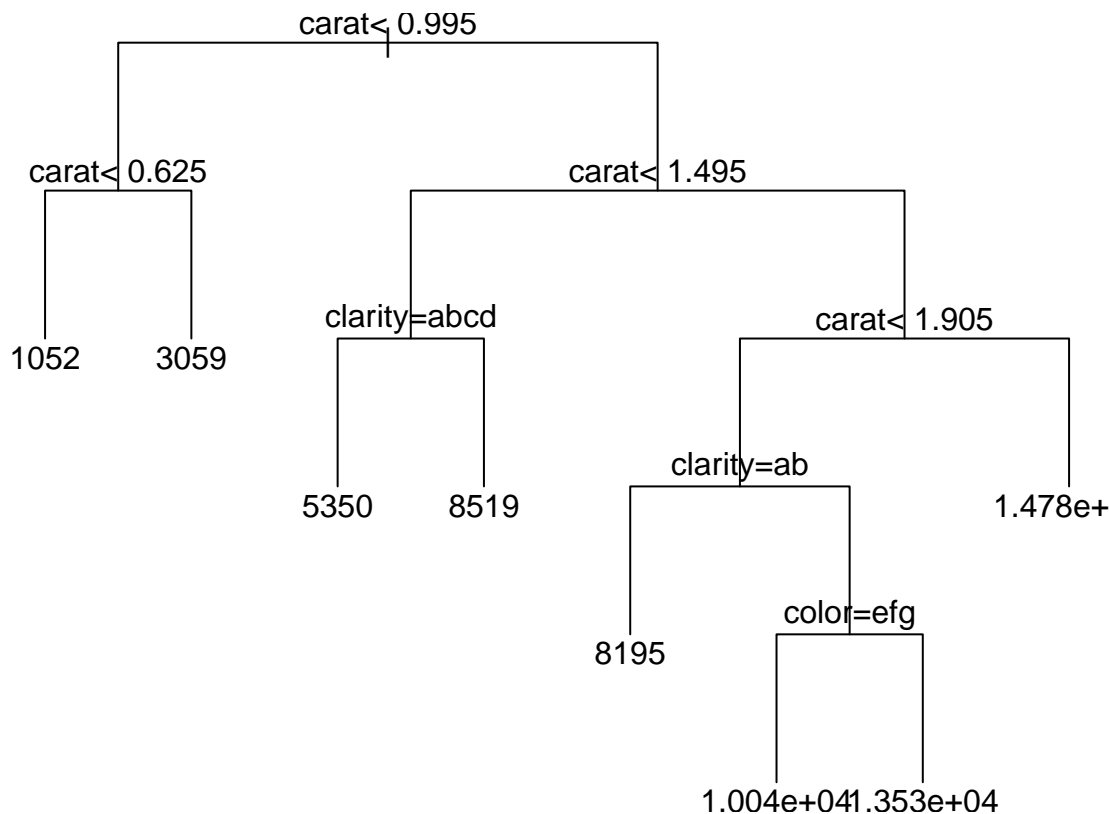
```
plot(fit, uniform=TRUE)
text(fit, cex=1)
```



... What if we remove the x, y, z, data, and table features? These above group are features that are not easily available to the common consumer, like me shopping for a new ring. Therefore, I'll see how the fit is with only the readily available features, carat, cut, color, and clarity.

```
# train the tree
fit2 <- rpart(price ~ carat + cut + color + clarity, data = splitData$train)
```

```
#plot the regression tree
plot(fit2, uniform=TRUE)
text(fit2, cex=1)
```

I'LL USE THE 'fit2' MODEL SINCE IT MAKES MORE SENSE FOR MY ENGAGEMENT RING SHOPPING SINCE CARAET, CLARITY, ETC ARE THINGS I CAN EASILY FIND OUT FROM A SELLAR, WHEREAS 'x, y, and z'ARE NOT.

DECISION TREE MODEL PREDICT DIAMOND PRICES Now, that I've generated a decision tree model, I'll now use it to predict prices.

DECISION TREE MODEL EVALUATION

MAE (Mean absolute Error)

```
maeTree <- mae(model = fit2, data=splitData$test)
maeTree
```

```
## [1] 835.146
```

I'll build a function to help compare MAE scores from different values for the tree depth (maxdepth)...

*# A function to get the maximum average error for a given max depth. You should pass in
the target as the name of the target colum and the predictors as vector where each
item in the vector is the name of the column.*

```
get_mae <- function(maxdepth, target, predictors, training_data, testing_data){
  #turn the predictors & target into a formula to pass to rpart
  predictors <- paste(predictors, collapse='+')
  formula <- as.formula(paste(target, "~", predictors, sep = ""))
  #build our model
  model <- rpart(formula, data=training_data, control = rpart.control(maxdepth = maxdepth))
  #get the mae
  mae <- mae(model, testing_data)
  return(mae)
}
```

```

#Feed in the target and predictors
target <- "price"
predictors <- c("carat", "cut", "color", "clarity")
# get the MaE for the maxdepths between 1 and 10
for(i in 1:10){
  mae <- get_mae(maxdepth = i, target = target, predictors = predictors,
                 training_data = splitData$train, testing_data = splitData$test)
  print(glue::glue("Maxdepth: ",i, "\t MAE: ", mae))
}

```

```

## Maxdepth: 1    MAE: 1717.50055497977
## Maxdepth: 2    MAE: 1042.80304167472
## Maxdepth: 3    MAE: 891.329899980287
## Maxdepth: 4    MAE: 865.382423661785
## Maxdepth: 5    MAE: 835.146013826217
## Maxdepth: 6    MAE: 835.146013826217
## Maxdepth: 7    MAE: 835.146013826217
## Maxdepth: 8    MAE: 835.146013826217
## Maxdepth: 9    MAE: 835.146013826217
## Maxdepth: 10   MAE: 835.146013826217

```

RANDOM FOREST

```

fitRandomForest <- randomForest(price ~ carat + cut + color + clarity, data=splitData$train)
maeForest <- mae(model = fitRandomForest, data=splitData$test)
maeForest

```

```
## [1] 899.502
```

```
fitRandomForest
```

```

##
## Call:
## randomForest(formula = price ~ carat + cut + color + clarity,      data = splitData$train)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 1
##
##              Mean of squared residuals: 1794744
##              % Var explained: 88.66

```

LINEAR MODEL

```

fitLinear <- lm(price ~ carat + cut + color + clarity, data=splitData$train)
maeLinear <- mae(model = fitLinear, data=splitData$test)
maeLinear

```

```
## [1] 808.8377
```

```
summary(fitLinear)
```

```

##
## Call:
## lm(formula = price ~ carat + cut + color + clarity, data = splitData$train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max

```

```
## -12621.1   -676.8   -195.9    461.5  10403.2
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -3707.097    16.598  -223.346 < 2e-16 ***
## carat       8877.777    14.342   619.006 < 2e-16 ***
## cut.L       704.988    24.039    29.327 < 2e-16 ***
## cut.Q      -327.297    21.178   -15.455 < 2e-16 ***
## cut.C       173.304    18.462     9.387 < 2e-16 ***
## cut^4        9.332    14.813     0.630  0.529
## color.L    -1895.714    21.060   -90.016 < 2e-16 ***
## color.Q    -614.831    19.177   -32.061 < 2e-16 ***
## color.C    -166.173    17.923    -9.271 < 2e-16 ***
## color^4     17.897    16.476     1.086  0.277
## color^5    -92.515    15.561    -5.945 2.78e-09 ***
## color^6    -58.799    14.159    -4.153 3.29e-05 ***
## clarity.L   4218.951    36.418   115.848 < 2e-16 ***
## clarity.Q  -1840.129    33.996   -54.128 < 2e-16 ***
## clarity.C    928.774    29.157    31.854 < 2e-16 ***
## clarity^4   -366.856    23.360   -15.704 < 2e-16 ***
## clarity^5    193.423    19.140    10.106 < 2e-16 ***
## clarity^6   -11.405    16.720    -0.682  0.495
## clarity^7    119.332    14.726     8.104 5.49e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1152 on 37740 degrees of freedom
## Multiple R-squared:  0.9162, Adjusted R-squared:  0.9161
## F-statistic: 2.291e+04 on 18 and 37740 DF,  p-value: < 2.2e-16
```

```
coef(fitLinear)
```

```
## (Intercept)      carat      cut.L      cut.Q      cut.C      cut^4
## -3707.096881  8877.777180  704.988263 -327.296838  173.303738   9.332101
##      color.L      color.Q      color.C      color^4      color^5      color^6
## -1895.713745 -614.831194 -166.173302  17.896622 -92.515227 -58.799452
##      clarity.L      clarity.Q      clarity.C      clarity^4      clarity^5      clarity^6
##  4218.950602 -1840.128996  928.774461 -366.856191  193.423086 -11.404789
##      clarity^7
##    119.332028
```

LOGISTIC REGRESSION

```
fitLogistic <- glm(price ~ carat + cut + color + clarity, data=splitData$train)
maeLogistic <- mae(model = fitLogistic, data=splitData$test)
maeLogistic
```

```
## [1] 808.8377
```

```
summary(fitLogistic)
```

```
##
## Call:
## glm(formula = price ~ carat + cut + color + clarity, data = splitData$train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```

```

## -12621.1    -676.8    -195.9    461.5    10403.2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3707.097    16.598 -223.346 < 2e-16 ***
## carat       8877.777    14.342  619.006 < 2e-16 ***
## cut.L        704.988    24.039   29.327 < 2e-16 ***
## cut.Q       -327.297    21.178  -15.455 < 2e-16 ***
## cut.C        173.304    18.462    9.387 < 2e-16 ***
## cut^4         9.332    14.813    0.630  0.529
## color.L     -1895.714    21.060  -90.016 < 2e-16 ***
## color.Q      -614.831    19.177  -32.061 < 2e-16 ***
## color.C     -166.173    17.923   -9.271 < 2e-16 ***
## color^4       17.897    16.476    1.086  0.277
## color^5      -92.515    15.561   -5.945 2.78e-09 ***
## color^6      -58.799    14.159   -4.153 3.29e-05 ***
## clarity.L    4218.951    36.418  115.848 < 2e-16 ***
## clarity.Q   -1840.129    33.996  -54.128 < 2e-16 ***
## clarity.C     928.774    29.157   31.854 < 2e-16 ***
## clarity^4    -366.856    23.360  -15.704 < 2e-16 ***
## clarity^5     193.423    19.140   10.106 < 2e-16 ***
## clarity^6    -11.405    16.720   -0.682  0.495
## clarity^7     119.332    14.726    8.104 5.49e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 1327262)
##
## Null deviance: 5.9747e+11  on 37758  degrees of freedom
## Residual deviance: 5.0091e+10  on 37740  degrees of freedom
## AIC: 639527
##
## Number of Fisher Scoring iterations: 2

REPORT GENERATION

print('REPORT')

## [1] "REPORT"

print("Decision Tree")

## [1] "Decision Tree"

print(maeTree)

## [1] 835.146

print("Random Forest")

## [1] "Random Forest"

print(maeForest)

## [1] 899.502

print("Linear Regression")

## [1] "Linear Regression"

```

```
print(maeLinear)
```

```
## [1] 808.8377
```

```
print("Logistic Regression")
```

```
## [1] "Logistic Regression"
```

```
print(maeLogistic)
```

```
## [1] 808.8377
```